

EC2 setup

Kirtikumar Shinde

Exported on 07/21/2020

Table of Contents

1	Conda Setup	4
2	Spark Installation	5
3	Spark configuration	6
4	User profile setup	7
5	QB python packages	8
6	Python libraries	9
7	Jupyter kernel setup	10
7.1	Code to be added:	10
7.2	Image of the changes:	11
8	Environment Test	12
9	Attachments	14

If the client is handover the docker, then following steps would be needed to setup conda environment and spark on the client environment.

1 Conda Setup

- Run the commands from the below script to download and setup anaconda environment.
- This would setup the conda environment for the particular user.
- Run the commands one by one.
- After the command is complete, run the `conda --version` to make sure anaconda is installed.

```
#!/bin/bash
set -e

## Configurations
# Anaconda installation prefix
PREFIX="$HOME/anaconda3"
# Path for conda binary
CONDA_BIN="$PREFIX/bin/conda"
# Path for python binary
CONDA_PY="$PREFIX/bin/python"

## Install Anaconda
# Download the latest version of Anaconda installation file
echo -e "\nDownloading..."
curl -f# -o "$HOME/install_anaconda.sh" https://repo.anaconda.com/archive/$(curl -sf https://
repo.anaconda.com/archive/ | grep -Eo "Anaconda3-20(1[8-9]|2[0-9])\.[0-9]{1,2}-Linux-x86_64\.sh" | head -1)

# Install Anaconda
echo -e "\nInstalling..."
bash "$HOME/install_anaconda.sh" -b -f -p $PREFIX
rm -f "$HOME/install_anaconda.sh"

## Config Anaconda
# Update conda to the latest version
$CONDA_BIN update -y conda

# Initialize conda for shell interaction
$CONDA_BIN init bash

# Set conda base env Python version to 3.7 then create the kernel for Jupyter
echo -e "\nSetting up Conda environment and Jupyter kernel..."
$CONDA_BIN install -y python=3.6
$CONDA_PY -m ipykernel install --user --name python3 --display-name "Python 3 (conda)"

echo -e "\nInstall Succeeded!"
```

2 Spark Installation

- Run the below commands to download the Spark. If you are looking for different version of Spark and Hadoop - change them accordingly.
- Make sure commands are run one by one.

```
##Spark Setup:
# make sure you are logged in as root to run below commands

export SPARK_VERSION=2.4.0
export HADOOP_VERSION=3.2.1

apt-get update -y
apt-get install default-jre -y
apt-get install default-jdk -y

mkdir /usr/local/spark/
chmod -R 755 /usr/local/spark/

cd /usr/local/spark/

curl https://archive.apache.org/dist/spark/spark-$SPARK_VERSION/spark-$SPARK_VERSION-bin-without-hadoop.tgz
--output spark-$SPARK_VERSION-bin-without-hadoop.tgz
tar -xvzf spark-$SPARK_VERSION-bin-without-hadoop.tgz

curl https://archive.apache.org/dist/hadoop/common/hadoop-$HADOOP_VERSION/hadoop-$HADOOP_VERSION.tar.gz --
output hadoop-$HADOOP_VERSION.tar.gz
tar -xvf hadoop-$HADOOP_VERSION.tar.gz

rm -rf spark-$SPARK_VERSION-bin-without-hadoop.tgz
rm -rf hadoop-$HADOOP_VERSION.tar.gz

ln -s /usr/local/spark/ /opt/spark
```

3 Spark configuration

- Edit the `spark-default.conf` and append the contents as below in the file.
- If the file does not exist, create this file from the template file available in the `/usr/local/spark/spark-2.4.0-bin-without-hadoop/conf/` directory.
- Make sure you only add the following lines at the end of the file. Please note - depending on the size of memory on EC2 you would have change the values for:
 - spark.driver.memory - should be set to around 70% of total RAM

```
spark.hadoop.fs.s3a.impl org.apache.hadoop.fs.s3a.S3AFileSystem
spark.hadoop.fs.s3a.connection.maximum 100
spark.jars.packages org.apache.hadoop:hadoop-aws:3.1.1
spark.driver.memory 180g
spark.driver.maxResultSize 10g
spark.driver.memoryOverhead 5g
spark.debug.maxToStringFields 10000
```

4 User profile setup

- Edit the user's .bashrc file using the command ``vi ~/.bashrc'` and add the following lines. Make sure these lines are added the end of the file without impacting any existing commands in the file.

```
export SPARK_VERSION=2.4.0
export HADOOP_VERSION=3.1.1
export PATH="$PATH:/opt/spark/spark-$SPARK_VERSION-bin-without-hadoop/bin"
export JAVA_HOME=/usr/lib/jvm/java-8-openjdk-amd64/jre
export SPARK_HOME=/opt/spark/spark-$SPARK_VERSION-bin-without-hadoop
export SPARK_DIST_CLASSPATH=/opt/spark/hadoop-$HADOOP_VERSION/etc/hadoop:/opt/spark/hadoop-$HADOOP_VERSION/share/hadoop/common/lib/*:/opt/spark/hadoop-$HADOOP_VERSION/share/hadoop/common/*:/opt/spark/hadoop-$HADOOP_VERSION/share/hadoop/hdfs:/opt/spark/hadoop-$HADOOP_VERSION/share/hadoop/hdfs/lib/*:/opt/spark/hadoop-$HADOOP_VERSION/share/hadoop/hdfs/*:/opt/spark/hadoop-$HADOOP_VERSION/share/hadoop/mapreduce/lib/*:/opt/spark/hadoop-$HADOOP_VERSION/share/hadoop/mapreduce/*:/opt/spark/hadoop-$HADOOP_VERSION/share/hadoop/yarn:/opt/spark/hadoop-$HADOOP_VERSION/share/hadoop/yarn/lib/*:/opt/spark/hadoop-$HADOOP_VERSION/share/hadoop/yarn/*
```

5 QB python packages

- On the home directory, create a folder as 'qb_python_packages'.
- Copy the provided QB python packages there. The directory should have following two packages:
 - PAI

6 Python libraries

- First activate the conda environment which you have created.
- Download the GIT repository for the project and go to that directory. Alternatively you can also download the list from page [1.2 Python package list](#)(see page 3) and add it to 'requirements.txt'.
- Before install libraries make sure the QB packages are copied manually to the home directory of user in the folder 'qb_python_packages'.
- *These steps would report some warnings/error about package versions. Please ignore these errors or warnings.*
- Assumptions:
 - It is 'base' conda environment.
 - The PAI package path is in the '/home/unix_user/qb_python_packages'.
 - The git repo is copied under the folder '/home/unix_user/repo/project_baldur'

```
# Activate conda environment
conda activate base

# Go to project directory
cd /home/unix_user/repo/project_baldur

# Run the python package installation command
pip install -r src/requirements.txt --find-links /home/unix_user/qb_python_packages
```

7 Jupyter kernel setup

- First find out the Jupyter kernel file name using the command 'jupyter kernelspec list'.
- Then edit that file to add following variables in it.
- *Note:*
 - Only add lines from 'env' onwards.
 - Code below show the final file.
 - Image below shows **only the text** which you have to add including the **comma**.
- Assumptions:
 - Kernel filename output from the command is '/home/unix_user/.local/share/jupyter/kernels/python3/kernel.json'

7.1 Code to be added:

```

1  {
2  "argv": [
3  "/home/unix_user/anaconda3/envs/test_new/bin/python",
4  "-m",
5  "ipykernel_launcher",
6  "-f",
7  "{connection_file}"
8  ],
9  "display_name": "new_env_python",
10 "language": "python",
11 "env": {
12   "SPARK_LOCAL_IP": "127.0.0.1",
13   "SPARK_VERSION": "2.4.0",
14   "HADOOP_VERSION": "3.1.1",
15   "PATH": "/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin:/snap/bin:/opt/spark/
16   spark-2.4.0-bin-without-hadoop/bin",
17   "SPARK_HOME": "/opt/spark/spark-2.4.0-bin-without-hadoop",
18   "SPARK_DIST_CLASSPATH": "/opt/spark/hadoop-3.1.1/etc/hadoop:/opt/spark/hadoop-3.1.1/share/hadoop/
19   common/lib/*:/opt/spark/hadoop-3.1.1/share/hadoop/common/*:/opt/spark/hadoop3.1.1/share/hadoop/
20   hdfs:/opt/spark/hadoop-3.1.1/share/hadoop/hdfs/lib/*:/opt/spark/hadoop-3.1.1/share/hadoop/hdfs/*:/
21   opt/spark/hadoop-3.1.1/share/hadoop/mapreduce/lib/*:/opt/spark/hadoop-3.1.1/share/hadoop/
22   mapreduce/*:/opt/spark/hadoop-3.1.1/share/hadoop/yarn:/opt/spark/hadoop-3.1.1/share/hadoop/yarn/
   lib/*:/opt/spark/hadoop-3.1.1/share/hadoop/yarn/*",
18   "PYSPARK_PYTHON": "/home/unix_user/anaconda3/bin/python",
19   "PYTHONSTARTUP": "/opt/spark/spark-2.4.0-bin-without-hadoop/python/pyspark/shell.py"
20   }
21 }
22 }
```

7.2 Image of the changes:

```
{
  "argv": [
    "/home/unix_user/anaconda3/envs/test_new/bin/python",
    "-m",
    "ipykernel_launcher",
    "-f",
    "{connection_file}"
  ],
  "display_name": "new_env_python",
  "language": "python",
  "env": {
    "SPARK_LOCAL_IP": "127.0.0.1",
    "SPARK_VERSION": "2.4.0",
    "HADOOP_VERSION": "3.1.1",
    "PATH": "/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin:/snap/bin:/opt/spark/spark-2.4.0-bin-without-hadoop/bin",
    "SPARK_HOME": "/opt/spark/spark-2.4.0-bin-without-hadoop",
    "SPARK_DIST_CLASSPATH": "/opt/spark/hadoop-3.1.1/etc/hadoop:/opt/spark/hadoop-3.1.1/share/hadoop/common/lib/*:/opt/spark/hadoop-3.1.1/share/hadoop/common/*:/opt/spark/hadoop3.1.1/share/hadoop/hdfs:/opt/spark/hadoop-3.1.1/share/hadoop/hdfs/lib/*:/opt/spark/hadoop-3.1.1/share/hadoop/hdfs/*:/opt/spark/hadoop-3.1.1/share/hadoop/mapreduce/lib/*:/opt/spark/hadoop-3.1.1/share/hadoop/mapreduce/*:/opt/spark/hadoop-3.1.1/share/hadoop/yarn:/opt/spark/hadoop-3.1.1/share/hadoop/yarn/lib/*:/opt/spark/hadoop-3.1.1/share/hadoop/yarn/*",
    "PYSPARK_PYTHON": "/home/unix_user/anaconda3/bin/python",
    "PYTHONSTARTUP": "/opt/spark/spark-2.4.0-bin-without-hadoop/python/pyspark/shell.py"
  }
}
```

8 Environment Test

- Test the environment to make sure everything is working.
- To do this open jupyter notebook. Make sure the same user for which above setting is done is logging to notebook.
- Make sure kernel is selected as the one created in step "Conda setup".
- In Jupyter, select 'shutdown kernel'. Give about 30 seconds and then select 'restart kernel'.
- This will make sure spark variables above are picked up by kernel.
- Then in the notebook write the below code. Make sure commands are written in different cells as marked in the code below, so each cell is run sequentially before other.
- The output from reading the data should show some data to confirm that environment is working fine.
- Assumptions:
 - S3 path of the data is: `s3://rwe-study-data/data/data/a_raw/lu_procedure/*`
 - Kernel in jupyter is selected - "Python 3 (conda)"

```

# Cell 1 - init spark
import findspark
findspark.init()

# Cell 2 - Create spark session
import pyspark
from pyspark.sql import SparkSession
spark = SparkSession.builder.master("local[*]").getOrCreate()

# Cell 3 - Check spark session is set
spark

# Cell 4 - import kedro
import pyspark.sql
from kedro.io import DataCatalog
from kedro.contrib.io.pyspark import SparkDataSet

# Cell 5 - Initialize file using kedro
spark_proc_ds = SparkDataSet(
    filepath="s3a://rwe-study-data/data/data/a_raw/lu_procedure/*",
    file_format="csv",
    load_args={"header": False, "inferSchema": False, "sep": "|",
    'schema':
        ''' category_dtl_cd          STRING,
category_dtl_code_desc  STRING,
category_genl_cd        STRING,
category_genl_code_desc STRING,
proc_cd                 STRING,
proc_desc               STRING,
proc_end_datex          STRING,
proc_typ_cd             STRING,
proc_end_date           LONG'''},
    save_args={"sep": "|", "header": True}
)
catalog = DataCatalog({'raw_lu_procedure' : spark_proc_ds})

# Cell 6 - read the data
df_proc = catalog.load('raw_lu_procedure')

# Cell 7 - Display the
df_proc.show()

# Cell 8 - count rows
df_proc.count()

```

9 Attachments



pai-0.18.0-py3-none-any.whl